


Enhanced Default Ergonomics

JVM “tuning” for beginners

Microsoft

Problem


Question 4 of 5: Do you trust JVM Ergonomics to pick the best GC for you?



- JVM tuning is hard, so developers trust the defaults
- Developers are deploying small JVMs, often in containers
 - Small memory footprint (often 0.5 to 2 GB)
 - Small CPU amount (often, “1 core / 1000 milli-cores”)
- Not setting max heap size
 - Often get default ergonomics (50% or 25%)
- Not selecting GC
 - Often, SerialGC (not bad)
 - But they expected the “*default GC*” - G1GC
 - Or worse, select a GC ignoring consequences of constrained resources


Some data

New Relic 2021




Some data

New Relic 2023




Some data

New Relic 2023




Assumption in 2000



Bare metal “shared”

- 2 Cores
- 2 GB of RAM
- Shared resources

Assumption in 2023



Container/VM “dedicated”

- 2 Cores
- 2 GB of RAM
- Isolated resources

Static JVM Ergonomics – Current state (as of 2000)

- Default heap size (of available memory):
Maximum

<256 MB	50%
256-512 MB	~126 MB
>512 MB	25%

Initial

<= 512 MB	8 MB
>=512 MB	1.5625%-1.7%

Minimum

64 MB – 8192 MB	8 MB
-----------------	------

- **GC selection**

Memory	Processors	GC selected
>=1792 MB	>1	G1
Otherwise		Serial


- **GC Threads**

GC	# of CPUs	ConcGCThreads	ParallelGCThreads
SerialGC	any	0	0
G1 GC	1-8	max((ParallelGCThreads+2)/4, 1)	ceil(#CPUS)
G1 GC	>8	max((ParallelGCThreads+2)/4, 1)	8 + (#CPUS-8) * (5/8)

- **Available processors**

- Server: 1 “active processor” for each CPU
- Container: cgroups-based counting on cpu_quota/cpu_period


Assumption in 2000



Bare metal
“shared”

- 2 Cores
- 2 GB of RAM
- Shared resources

Assumption in 2023



Container/VM
“dedicated”

- 2 Cores
- 2 GB of RAM
- Isolated resources

Static JVM Ergonomics – Container-aware (2023)

- Default heap size (of available memory):
Maximum



Initial



Minimum



- **GC selection**



- **GCThreads**



- **Available processors**



Ideas

- Produce generic Java runtimes for containers with tweaked defaults
- Add more if/elses into `select_gc_ergonomically()` `set_heap_size()`
- JVM launcher configuration files
- Stuff we haven't thought of 😊
- Ergonomics profiles

Naïve and simplistic proposal

- Ergonomics Profile
 - shared: traditional ergonomics (unchanged)
 - dedicated: when the JVM believes it will be likely alone (e.g. containers)

Memory	Processors	GC selected
Any	1	Serial
<=2048 MB	>1	ParallelGC
>2048 MB	>1	G1
>=16 GB	>1	ZGC

Memory	Heap size
< 0.5 GB	50%
>= 0.5 GB	75%
>= 4 GB	80%
>= 6 GB	85%
>= 16 GB	90%

Should we disable `-XX:-UseDynamicNumberOfCompilerThreads` for single proc and <=512MB envs ??

Works on VMs too (image templates)

```
JAVA_TOOL_OPTIONS="-XX:ErgonomicsProfile=dedicated"
```

Challenges

- Sooner
 - [JDK-8261242](#): `is_containerized()` returns true when run outside a container
 - [JDK-8134507](#): Improve the tiered compilation compiler thread ergonomics
 - Conscious of non-heap consumption: code cache, direct memory, metaspace, GC, etc.
- Later
 - [JDK-8254091](#): Need a mechanism (and API) to reliably determine if a JVM is executing in a container context
 - What to do with Client/Server Class Machines feature?

Diving into code

- <https://github.com/microsoft/openjdk-jdk/pull/9>

Next steps

- Confirm settings for “dedicated” profile are fair. Any feedback?
 - What other flags should be enabled, and under what conditions?
 - What other profiles may come in the future? Hardware-based profiles?
 - How does this ties to Leyden profiles?
 - Is this better than status quo?
-
- Is this a good path forward to improve new Java app onboarding in the Cloud?